

Multi-threading on CPUs with OpenMP and Metrics for Performance Analysis of Applications

Polytech Sorbonne – EI5-SE – Calcul haute performance (EPU-F9-IHP)

Adrien CASSAGNE

January 5, 2026



Table of Contents

1 Introduction to OpenMP

- ▶ Introduction to OpenMP
- ▶ OpenMP Use Cases
- ▶ Parallel Code Analysis
- ▶ Kernel Performance Analysis
- ▶ References



Programming Multi-core CPUs

1 Introduction to OpenMP

- Nowadays, multi-core architecture is well spread in High Performance Computing (HPC) and in embedded targets
- There are two main ways to use multi-core architectures
 1. Create **multiple processes** (= distributed memory model)
→ MPI standard, Unix inter-processes communications, sockets, ...
 2. Create **multiple threads** (= shared memory model)
→ Threads POSIX, OpenMP, ...
- In this session we will not talk about the multiple processes model
- And **we will go deeper into the multi-threaded model**



OpenMP Presentation

1 Introduction to OpenMP

- OpenMP is a language dedicated to setup multi-threaded codes
- It is based on **compiler directives** (`#pragma`)
 - Those directives describe how to perform the parallelism
 - The main advantage of directives is to **not modify sequential code** (in theory...)

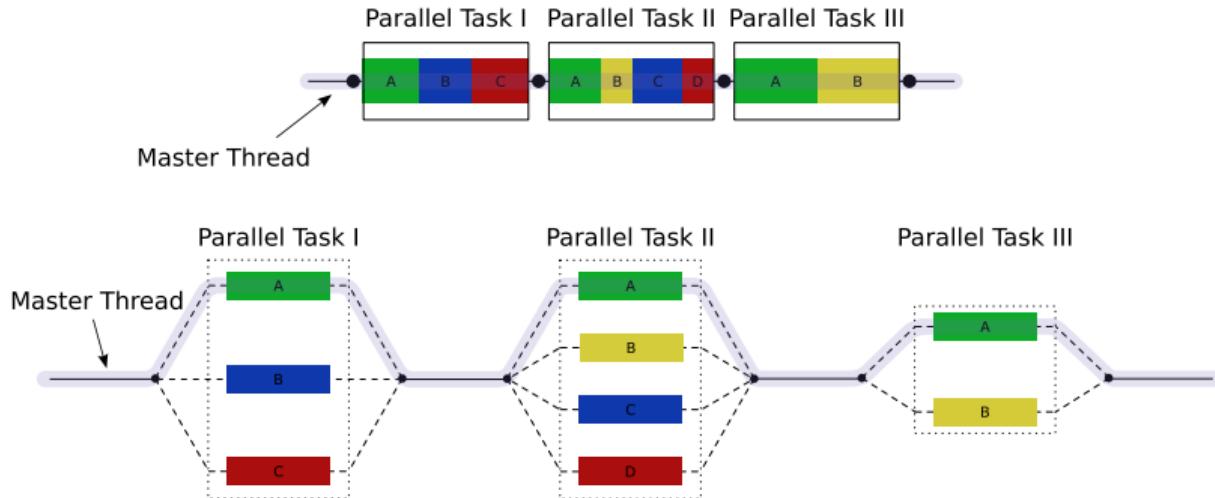
```
1 void add_vectors(const float* A, const float* B, float* C, const size_t n)
2 {
3 #pragma omp parallel // directive for the creation of a parallel zone (= threads creation)
4 { // <- beginning of the parallel zone
5 #pragma omp for // directive for distribution of for-loop indices among threads
6   for (size_t i = 0; i < n; i++)
7     C[i] = A[i] + B[i];
8 } // <- end of the parallel zone
9 }
```

Simple `add_vectors` OpenMP implementation



Fork-join Model

1 Introduction to OpenMP

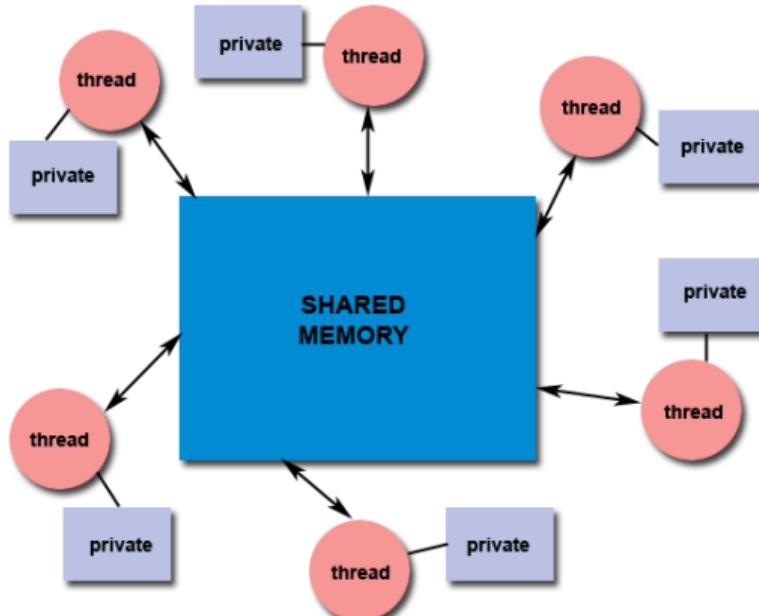


- When using `#pragma omp parallel` directive: threads are **created** (= **fork**)
- At the end of a parallel zone
 - Threads are **destroyed** (= **join**), except for the master thread
 - There is an implicit barrier



Shared Memory Model

1 Introduction to OpenMP

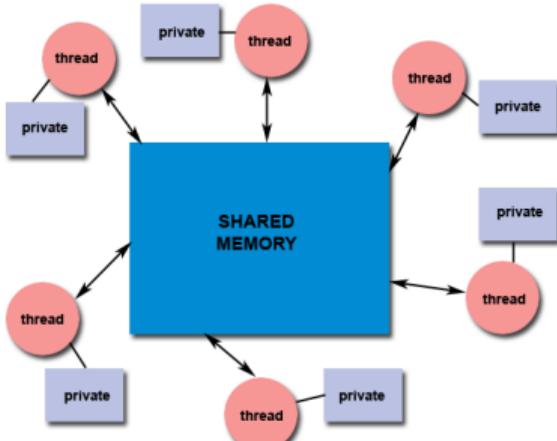


- Each thread can access the global memory zone
 - This is called the shared memory (or the RAM of the CPUs)
- But threads also own private data
 - Not completely shared model
 - Very often, the key for achieving performance is to keep the memory private when possible...



Shared Memory Model – Code Example

1 Introduction to OpenMP



```
1 void add_vectors(const float* A, const float* B, float* C,
2                   const size_t n)
3 {
4 #pragma omp parallel
5 {
6 #pragma omp for
7     // 'i' is private because it is declared
8     // after the omp parallel directive
9     for (size_t i = 0; i < n; i++) // <- 'n' is shared
10        // 'A', 'B' and 'C' are shared!
11        C[i] = A[i] + B[i];
12    }
13 }
```

- By default, variables that are **declared before a parallel zone** are **shared** (here A, B, C and n)
- And variables **declared inside a parallel zone** are **private** (here i)



Control Data Range

1 Introduction to OpenMP

- OpenMP provides data range control
 - **private**: local to the thread,
 - **firstprivate**: local to the thread and initialized
 - **shared**: shared by all the threads, in C/C++ this is the default behavior
- Here **alpha** is a constant, we can put it in the private memory of each thread
- Efficient parallelism comes with minimal synchronizations
 - Shared data can generate a lot of synchronizations
 - Privacy increases thread independence

```
1 void dot(const float* A,
           float* B,
           const float alpha,
           const size_t n)
2 {
3     #pragma omp parallel    \
4         shared(A, B)    \
5         firstprivate(alpha, n)
6     {
7         #pragma omp for
8             // 'i' is still private because
9             // it is declared after the
10            // parallel zone
11            for (size_t i = 0; i < n; i++) {
12                B[i] = alpha * A[i];
13            }
14        }
15    }
16 }
```



for-loop Indices Distribution (1)

1 Introduction to OpenMP

- for-loop indices distribution can be controlled by the **schedule** clause
 - **static**: indices distribution is precomputed (at compilation time), and the amount of indices is the same for each thread
 - **dynamic**: indices distribution is done in real time along the loop execution, work load balancing can be better than with the **static** scheduling but dynamic scheduling costs some additional resources in order to attribute indices at real time. This is a **greedy algorithm**.
- There are other types of scheduling and it is also possible to choose the scheduler at runtime with **OMP_SCHEDULE** environment variable and the **schedule(runtime)** primitive

```
1 // ...
2 #pragma omp for schedule(static, 128) // we statistically attribute 128 per 128 indices to each threads
3     for(int i = 0; i < n; i++)
4         B[i] = alpha * A[i];
5 // ...
```



for-loop Indices Distribution (2)

1 Introduction to OpenMP

Static



Dynamic



Guided



→
Iteration number

- Thread 0
- Thread 1
- Thread 2

From DOI: 10.1007/s10291-022-01266-8.



Sections

1 Introduction to OpenMP

- OpenMP proposes **sections** to enable threads to execute different “parts” (= **section**) of the code

```
1 #pragma omp parallel
2 { // <- multiple threads have been created
3 #pragma omp sections
4 { // <- entering in a 'sections' zone
5 #pragma omp section
6 { // this code is executed by a single thread
7     printf("id = %d,", omp_get_thread_num());
8 }
9 #pragma omp section
10 { // this code is executed by a single thread too
11     printf("id = %d,", omp_get_thread_num());
12 }
13 } // <- end of the 'sections' zone
14 } // <- threads are destroyed
```

- The code will print: id = 0, id = 1,
- But not necessarily in this order



Single and master directives

1 Introduction to OpenMP

```
1 void dot(const float* A, float* B, const float alpha,
2           const size_t n) {
3     #pragma omp parallel shared(A, B) \
4                     firstprivate(alpha, n)
5     {
6         #pragma omp for
7         for (size_t i = 0; i < n; i++) {
8             B[i] = alpha * A[i];
9         #pragma omp single
10        { // executed by only one thread
11            printf("B[i] = %f," B[i]);
12        } // <- there is an implicit barrier here, all the
13          // threads are waiting
14    }
15 }
16 }
```

```
1 void dot(const float* A, float* B, const float alpha,
2           const size_t n) {
3     #pragma omp parallel shared(A, B) \
4                     firstprivate(alpha, n)
5     {
6         #pragma omp for
7         for (size_t i = 0; i < n; i++) {
8             B[i] = alpha * A[i];
9         #pragma omp master
10        { // executed by only the master thread
11            printf("B[i] = %f," B[i]);
12        } // <- there is NO implicit barrier here, other
13          // threads will not wait for the master thread
14    }
15 }
16 }
```

- Directives to be executed by a single thread
- Even if they look similar, there is an implicit barrier after the `single` directive



Environment Variables

1 Introduction to OpenMP

- OpenMP take advantage of environment variable to customize the multi-threaded execution
- The most famous one is `OMP_NUM_THREADS`, it enables control the number of threads in the OpenMP parallel zones

```
1 export OMP_NUM_THREADS=4
2 ./my_omp_program # the code will run on 4 threads
3
4 OMP_NUM_THREADS=2 ./my_omp_program # the code will run on 2 threads
```

- It is possible to select which scheduler will be used with the `OMP_SCHEDULE` environment variable

```
1 export OMP_NUM_THREADS=4
2 OMP_SCHEDULE="static,3" ./my_omp_program # the code will run on 4 threads and
                                             # with a static scheduling of size 3
3
```



And Many more Features

1 Introduction to OpenMP

- Tasks
- Atomic, critical section
- SIMD
- Accelerators (GPUs)
- ...



Go Further

1 Introduction to OpenMP

- Previous slides were a brief overview of the main OpenMP principles
- To have more precise informations you can take a look at the very good OpenMP reference card¹
 - It could be a very good idea to print it and keep it ;-)
- In the next slides we will pay attention to some **OpenMP use cases**

¹<https://www.openmp.org/resources/refguides/>



Table of Contents

2 OpenMP Use Cases

- ▶ Introduction to OpenMP
- ▶ OpenMP Use Cases
- ▶ Parallel Code Analysis
- ▶ Kernel Performance Analysis
- ▶ References



Avoid False Sharing

2 OpenMP Use Cases

- **False sharing** is a phenomena that occurs when threads write simultaneously data in a same cache line
 - Remember, the cache system works with lines of words: a line is the smallest element in caches coherence mechanism
 - If two or more threads are working on the same line they cannot write data simultaneously!
→ Stores are serialized and we talk about false sharing
- To avoid false sharing, threads have to work on a **bigger amount of data than the cache line size**
 - Concretely we have to avoid `(static,1)` or `(dynamic,1)` scheduling
 - Cache lines are not very big (≈ 64 Bytes)
 - Just putting a `(static,16)` or `(dynamic,16)` often resolves the problem
 - Be aware that in some OpenMP implementations, the default scheduling pattern is `(static,1)`!



Threads Synchronizations – Barriers

2 OpenMP Use Cases

- In OpenMP there are a lot of **implicit barriers**, after each
 - `#pragma omp parallel` directive
 - `#pragma omp for` directive
 - `#pragma omp single` directive
- But not after `#pragma omp master` directive!
- If we are sure that there is no need to synchronise threads after the `#pragma omp for` directive, we can use the **nowait** clause
- Optimally we need only one `#pragma omp parallel` directive in a fully parallel code
 - OpenMP manages a pool of threads in order to reduce the cost of the `#pragma omp parallel` directive but this is not free, each time OpenMP has to reorganize the pool and wakes up the required threads



Threads Synchronizations – Barriers – Example

2 OpenMP Use Cases

```
1 // A, B & C <- size = n, D <- size = 2n
2 void kernel_v1(const float *A, const float *B, const float *C,
3                 float *D, const float alpha, const size_t n) {
4 // overhead: threads creation and private variables creation
5 #pragma omp parallel shared(A, B, D) \
6                         firstprivate(alpha, n)
7 {
8 #pragma omp for schedule(static,16)
9 {
10    for (size_t i = 0; i < n; i++)
11        D[i] = alpha * A[i] + B[i];
12 } // implicit barrier
13 } // implicit barrier
14
15 // overhead: threads attribution and private variables creation
16 #pragma omp parallel shared(A, C, D) firstprivate(n)
17 {
18 #pragma omp for schedule(static,16)
19 {
20    for (size_t i = 0; i < n; i++)
21        D[n + i] = A[i] + C[i];
22 } // implicit barrier
23 } // implicit barrier
24 }
```

```
1 // A, B & C <- size = n, D <- size = 2n
2 void kernel_v2(const float *A, const float *B, const float *C,
3                 float *D, const float alpha, const size_t n) {
4 // overhead: threads creation and private variables creation
5 #pragma omp parallel shared(A, B, C, D) \
6                         firstprivate(alpha, n)
7 {
8 #pragma omp for schedule(static,16) nowait
9 {
10    for (size_t i = 0; i < n; i++)
11        D[i] = alpha * A[i] + B[i];
12 } // no implicit barrier (nowait clause)
13
14 #pragma omp for schedule(static,16)
15 {
16    for (size_t i = 0; i < n; i++)
17        D[n + i] = A[i] + C[i];
18 } // implicit barrier
19 } // implicit barrier
20 }
21
22 /* 'kernel_v2' is a better version than 'kernel_v1':
23 *   - only one parallel zone
24 *   - no barrier after the first loop (nowait clause) */
```



Threads Synchronizations – Critical Sections

2 OpenMP Use Cases

- Sometimes it is not possible to have a fully parallel code and some regions of the code remain intrinsically sequential
- In OpenMP we can specify this kind of region with the **#pragma omp critical** directive
 - In a critical section, all the threads will execute the code but the execution is made one by one
 - Same as mutual exclusion (mutex) zones in POSIX threads
- But we have to use this directive carefully
 - **#pragma omp critical** can be the main cause of slow down in OpenMP codes!



Threads Synchros – Critical Sections – Example

2 OpenMP Use Cases

Scale A in B and find the minimum value of B in `min_val`

```
1 float kernel_v1(const float *A, float *B, const size_t n) {
2     float min_val = INF;
3     #pragma omp parallel shared(A, B, min_val) firstprivate(n)
4     {
5         #pragma omp for schedule(static,16)
6         {
7             for (size_t i = 0; i < n; i++) {
8                 B[i] = 0.5f * A[i];
9                 #pragma omp critical // we want to be sure that only one
10                // thread can modify min_val
11                 if (B[i] < min_val)
12                     min_val = B[i];
13             }
14         }
15     }
16 }
17     return min_val;
18 }
19
20 /* This code is slow because each loop step contains a
21 * sequential part */
```

```
1 float kernel_v2(const float *A, float *B, const size_t n) {
2     float min_val = INF;
3     #pragma omp parallel shared(A, B, min_val) firstprivate(n)
4     {
5         #pragma omp for schedule(static,16)
6         {
7             for (size_t i = 0; i < n; i++) {
8                 B[i] = 0.5f * A[i];
9                 // no more threads synchro to perform the test
10                 if (B[i] < min_val)
11                     #pragma omp critical
12                     {
13                         // this is very important to re-do the test because
14                         // an other thread may have modify the min_val value
15                         if (B[i] < min_val)
16                             min_val = B[i];
17                     }
18                 }
19             }
20     return min_val;
21 }
```



Search Algorithms

2 OpenMP Use Cases

- In OpenMP 3 there is no optimal solution for search algorithms
- This kind of algorithm typically requires while-loops or do-while-loops
- However there is a tip to fix this lack in OpenMP 3 (see next slide)
- Latest versions of OpenMP (v4 to v6) provides better control of threads
 - We can terminate threads...
 - But this lead to more complex solutions, we will not see them today



Search Algorithms – OpenMP 3 Tip

2 OpenMP Use Cases

Search if `val` element is in the `A` array of size `n`

```
1 bool search_val_v1(const float *A, const size_t n, float val)
2 {
3     bool found = false;
4     #pragma omp parallel shared(A, found) firstprivate(val)
5     {
6         #pragma omp for schedule(static,16)
7         {
8             for (size_t i = 0; i < n; i++) {
9                 if (A[i] == val)
10                     found = true; // no more break, this is valid but
11                         // this is also slow, no more early
12                         // exit :-(

13     }
14 }
15 }
16 return found;
17 }
```

```
1 bool search_val_v2(const float *A, const size_t n, float val)
2 {
3     bool found = false;
4     #pragma omp parallel shared(A, found) firstprivate(val)
5     {
6         #pragma omp for schedule(static,16)
7         {
8             for (size_t i = 0; i < n; i++) {
9                 if (!found) // we are doing nothing if we have found
10                     // the value in the array
11                     if (A[i] == val)
12                         found = true;
13         }
14     }
15 }
16 return found;
17 }
```



Table of Contents

3 Parallel Code Analysis

- ▶ Introduction to OpenMP
- ▶ OpenMP Use Cases
- ▶ Parallel Code Analysis
- ▶ Kernel Performance Analysis
- ▶ References



Execution Time

3 Parallel Code Analysis

How do you compare two versions of a code that does the same thing (from the functional point of view)?

- Compare the execution time of the two versions
 - The fastest program is the most efficient one
 - Intuitive and worth keeping in mind
- Be careful to compare the same times
 - **Classic error:** comparing the total execution time of one program with just a sub-part of another program's execution time
 - ➔ In this case, the two measured times are not comparable



Execution Time of a Parallel Code

3 Parallel Code Analysis

- Let's consider \mathcal{D}_1 (or \mathcal{D}_s) the sequential time (time on 1 core) of a code
 - With 2 cores, we can hope to divide the time by 2 at best ($\mathcal{D}_2^m \geq \mathcal{D}_s/2$)
 - With 3 cores, we can hope to divide the time by 3 at best ($\mathcal{D}_3^m \geq \mathcal{D}_s/3$)
- The following table shows the execution times measured for a Code 1:

# core (\mathcal{C})	Measured time (\mathcal{D}^m)	Optimal time (\mathcal{D}^o)
1	98 ms	98.0 ms
2	50 ms	49.0 ms
3	35 ms	32.7 ms
4	27 ms	24.5 ms
5	22 ms	19.6 ms
6	18 ms	16.3 ms

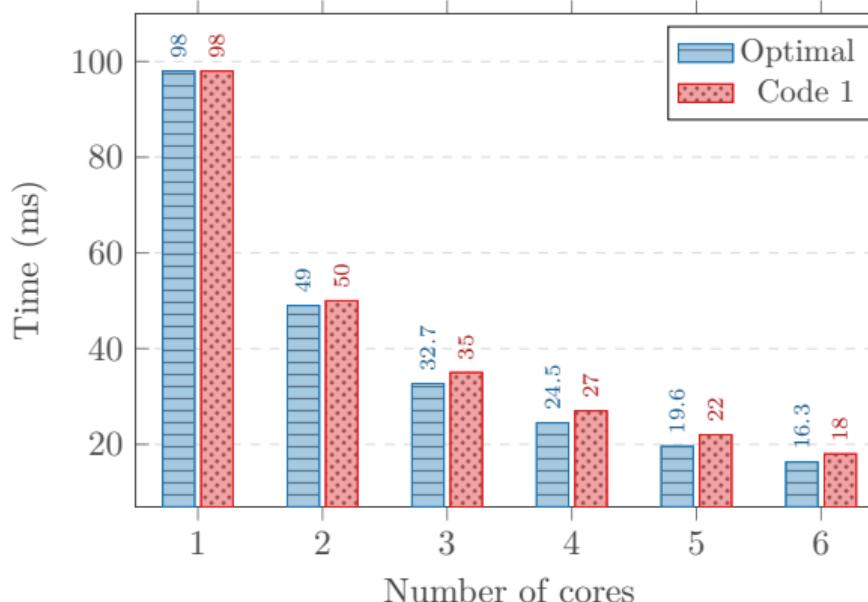
- Optimal time = $\mathcal{D}^o = \mathcal{D}_s/\mathcal{C}$



Visualization of the Execution Time

3 Parallel Code Analysis

- The previous table is not easy to read
- Let's look at the results on a graph:





Speedup – Definition

3 Parallel Code Analysis

$$\mathcal{S} = \mathcal{D}_s / \mathcal{D}_{\mathcal{C}},$$

with \mathcal{D}_s the time measured for the 1-core version (= sequential version) of the code and $\mathcal{D}_{\mathcal{C}}$ the time measured for the parallel version with \mathcal{C} cores.

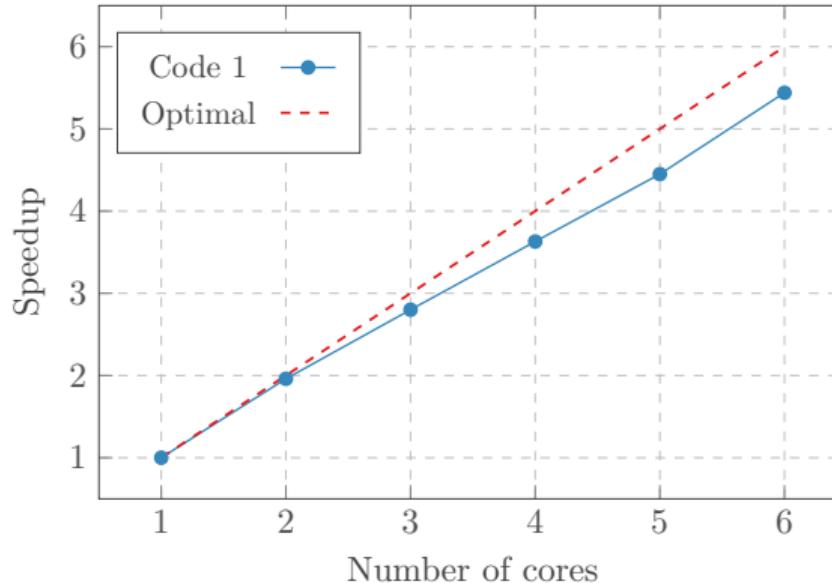
# cores (\mathcal{C})	Time (\mathcal{D}^m)	Speedup (\mathcal{S})
1	98 ms	1.00
2	50 ms	1.96
3	35 ms	2.80
4	27 ms	3.63
5	22 ms	4.45
6	18 ms	5.44

- Sequential time is used as reference time



Speedup – Visualization

3 Parallel Code Analysis



- Visually very simple to see if Code 1 is close to the the optimal
- Optimal speedup is equal to the number of cores used (no more!)



Amdahl's Law

3 Parallel Code Analysis

- Can we increase parallelism indefinitely to speedup our codes?
 - Amdahl said no!
 - To be more precise, it depends on the characteristics of the code...
 - If the code is fully parallelizable: speedup is infinite
 - If the code is NOT fully parallelizable: there's a limit

$$S_{\max} = \frac{1}{1 - f\mathcal{D}_p},$$

with S_{\max} the maximum achievable speedup and $f\mathcal{D}_p$ the fraction of parallel time in the code ($0 \leq f\mathcal{D}_p \leq 1$).



Amdahl's Law – Example

3 Parallel Code Analysis

- Let's take a code composed of two parts:
 - 20 % is intrinsically sequential
 - 80 % can be parallelized
- What is the maximum achievable speedup?

$$S_{\max} = \frac{1}{1 - f\mathcal{D}_p} = \frac{1}{1 - 0.8} = \frac{1}{0.2} = 5.$$

That's not much when you consider architectures with dozens of CPU cores!



Efficiency – Definition

3 Parallel Code Analysis

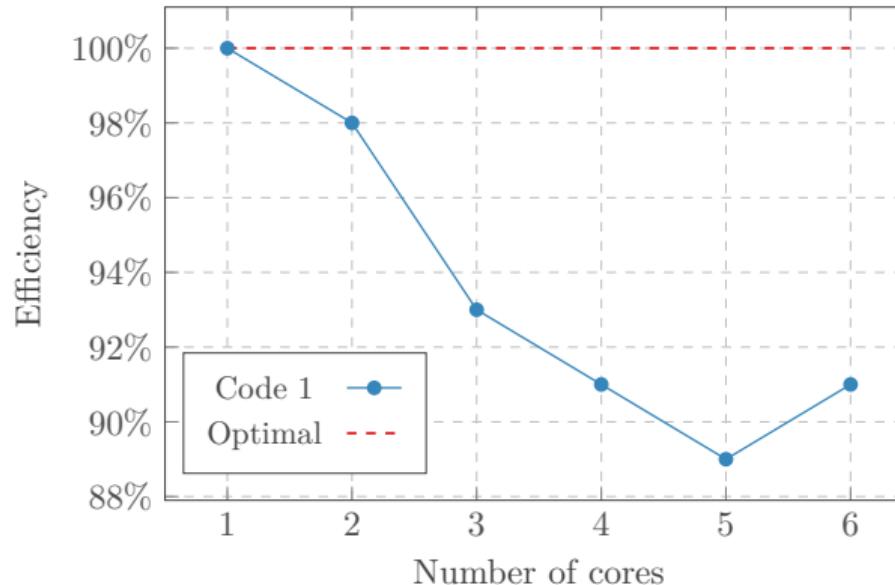
- Several ways to define the efficiency (\mathcal{E}) of a code
 - From the speedup: $\mathcal{E} = \mathcal{S}^m / \mathcal{S}^o$
 - From the execution time: $\mathcal{E} = \mathcal{D}^m / \mathcal{D}^o$
- Efficiency is a ratio: $0\% < \mathcal{E} \leq 100\%$
- For Code 1, the efficiency as a function of the number of cores:

# cores (\mathcal{C})	Time (\mathcal{D}^m)	Speedup (\mathcal{S})	Efficiency (\mathcal{E})
1	98 ms	1.00	100%
2	50 ms	1.96	98%
3	35 ms	2.80	93%
4	27 ms	3.63	91%
5	22 ms	4.45	89%
6	18 ms	5.44	91%



Efficiency – Visualization

3 Parallel Code Analysis



- Equivalent to the speedup, at least for now...



Scalability of a Code – Definition

3 Parallel Code Analysis

- The scalability of a code is its capacity to be efficient when the number of cores increases
 - A code scales if it is able to benefit from the power of several cores
- How do we measure a code's scalability? How do we know if a code doesn't scale?
 - No simple answer
- 2 widely used models for characterizing the scalability of parallel code:
 - The “strong” scalability
 - The “weak” scalability



Strong Scalability – Code 1 Example

3 Parallel Code Analysis

- Measures execution time as a function of the number of cores
- With a **constant problem size**
- For example, for **Code 1**, with a problem of size 100:

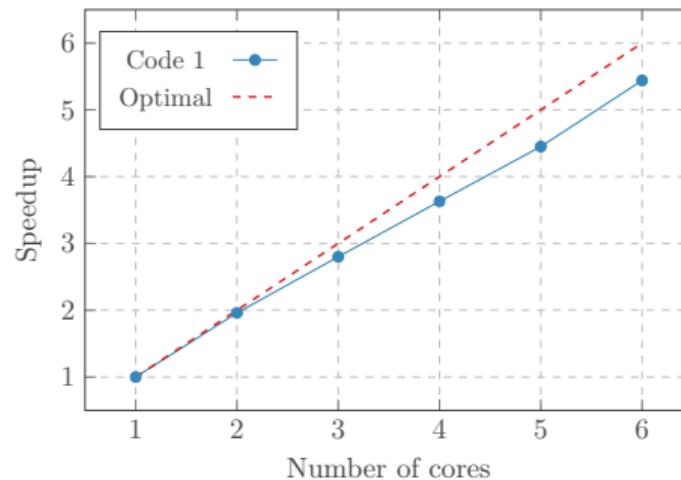
# cores	Size	Time	Speedup
1	100	98 ms	1.00
2	100	50 ms	1.96
3	100	35 ms	2.80
4	100	27 ms	3.63
5	100	22 ms	4.45
6	100	18 ms	5.44



Strong Scalability – Code 1 Visualization

3 Parallel Code Analysis

Strong scalability is generally observed on an speedup graph:



Here, for 6 cores, Code 1 achieves a speedup of 5.4, so we can conclude that this code scales well up to 6 cores.



Strong Scalability – Code 2 Example

3 Parallel Code Analysis

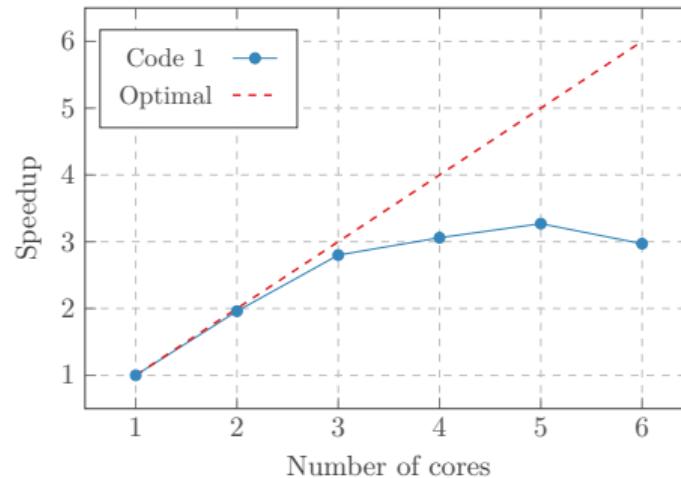
- Let's consider a new Code 2
- Here are the measurements for this code:

# cores	Size	Time	Speedup
1	100	98 ms	1.00
2	100	50 ms	1.96
3	100	35 ms	2.80
4	100	32 ms	3.06
5	100	30 ms	3.27
6	100	33 ms	2.97



Strong Scalability – Code 2 Visualization

3 Parallel Code Analysis



- We can see that the strong scalability of Code 2 is poor
- Above a certain number of cores, parallelism can no longer speedup code :-)



Weak Scalability

3 Parallel Code Analysis

- This model considers the execution time as a function of the number of cores
- **And the problem size increases in proportion to the number of cores!**
- Compute the *speedup* makes no sense if the problem size is not constant
- BUT it is possible to compute the efficiency: $\mathcal{E} = \mathcal{D}^s / \mathcal{D}^m$

Intuition: if we can't compute a problem of a given size any faster, can we compute a bigger problem in the same time?

Most of the time, yes, and it's easier! This is what happens most of the time in high performance computing: scientific models become more and more refined = the size of the problem increases.

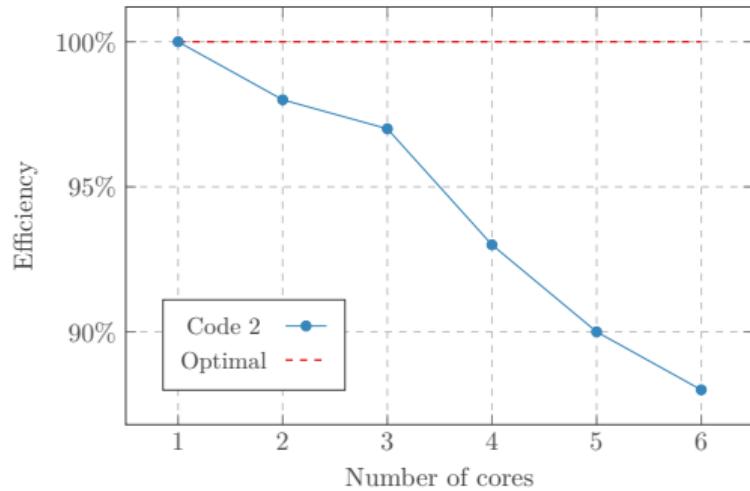


Weak Scalability – Code 2 Example

3 Parallel Code Analysis

Measures for Code 2:

# cores	Size	Time	Efficiency
1	100	098 ms	100%
2	200	100 ms	98%
3	300	101 ms	97%
4	400	105 ms	93%
5	500	109 ms	90%
6	600	111 ms	88%



- The weak scalability of Code 2 is good ($\approx 90\%$ for 6 cores)
- Why is strong scalability bad?
 - Amdahl's law: not enough parallelism for a small problem size

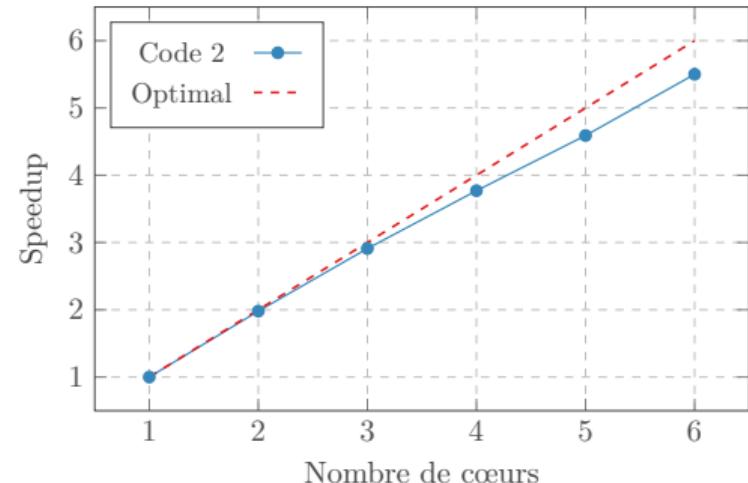


Strong Scalability – Code 2 AGAIN!

3 Parallel Code Analysis

- New problem size: 600

# cores	Size	Time	Speedup
1	600	611 ms	1.00
2	600	308 ms	1.98
3	600	210 ms	2.91
4	600	162 ms	3.77
5	600	133 ms	4.59
6	600	111 ms	5.50



- For larger problem sizes, strong scalability is good
- Not always possible to test for strong scalability: lack of time
- Not always possible to test for weak scalability: impossible to have ever larger problem sizes



Table of Contents

4 Kernel Performance Analysis

- ▶ Introduction to OpenMP
- ▶ OpenMP Use Cases
- ▶ Parallel Code Analysis
- ▶ Kernel Performance Analysis
- ▶ References



Number of Arithmetic Operations

4 Kernel Performance Analysis

- The number of arithmetic operations in a code is an **important characteristic**
- Example of the number of float operations (*flops*) in a kernel that performs a sum:

```
1 float sum(float *values, size_t n) {
2     float sum = 0.f;
3
4     // total flops = n * 1
5     for (size_t i = 0; i < n; i++)
6         sum = sum + values[i]; // 1 flop because of the addition
7
8     return sum;
9 }
```



Number of Operations per Second

4 Kernel Performance Analysis

- Metric widely used in high performance computing, particularly the number of **floating-point** operation per second (flop/s)
 - The same is defined for integer operation (iop/s)
 - Or simply the number of operation per second (op/s), or *Million Instructions Per Second* (MIPS)
 - This metric is used in the kernel Linux (`cat /proc/cpuinfo`)
- The flop/s ratio can be directly compared with the peak performance of a computing architecture
- A metric to estimate the good use (or not) of the hardware architecture



Processor Peak Performance

4 Kernel Performance Analysis

- **The processor's maximum computational capacity**
- It can be deduced from our knowledge of the hardware architecture:

$$\text{peakPerf} = nOps \times freq \times nCores,$$

with $nOps$ the number of operations the architecture can achieve in one cycle (ILP), $freq$ the processor frequency and $nCores$ the number of processor cores.



Processor Peak Performance – Example

4 Kernel Performance Analysis

CPU name	Core i7-2630QM
Architecture	Sandy Bridge
Vect. inst.	AVX-256 bit (4 double, 8 single)
Frequency	2 GHz
Nb. cores	4

Peak performance in floating-point single precision:

$$\text{peakPerf}_{sp} = nOps \times freq \times nCores = (2 \times 8) \times 2 \times 4 = 128 \text{ Gflop/s}$$

Peak performance in floating-point double precision:

$$\text{peakPerf}_{dp} = nOps \times freq \times nCores = (2 \times 4) \times 2 \times 4 = 64 \text{ Gflop/s}$$

- $nOps = 2 \times \text{vectorSize}$ because the architecture back-end allows to issue 2 instructions per cycle (vadd et vmul)



Arithmetic Intensity

4 Kernel Performance Analysis

- Sometimes (even often) measured op/s are far from peak performance
 - Code is poorly optimized
 - Peak performance cannot be achieved
 - In most cases, both are true...
- With **arithmetic intensity** we take into account **memory accesses**:

$$AI = \frac{ops}{memops}.$$



Arithmetic Intensity – Example

4 Kernel Performance Analysis

```
1 float sum(float *values, size_t n) {
2     float sum = 0.f; // do not count the memory acces to 'sum' because it will be optimized in register
3     // total flops = n * 1 // total memops = n * 1
4     for (size_t i = 0; i < n; i++)
5         sum = sum + values[i]; // 1 flop because of the addition, 1 memop because of one acces in
6                             // the 'values' array
7     return sum;
8 }
```

- Arithmetic intensity of `sum` is: $AI_{\text{sum}} = \frac{n \times 1}{n \times 1} = 1$
- The higher the arithmetic intensity, the more the code is limited by computational units
- The lower the arithmetic intensity, the more the code is limited by memory accesses



Operational Intensity

4 Kernel Performance Analysis

- Compared to arithmetic intensity, operational intensity takes into account **the size of the data in memory**:

$$OI = \frac{flops}{memops \times sizeOfData} = \frac{AI}{sizeOfData},$$

sizeOfData depends on the datatype, `int` and `float` use 4 bytes, `long long int` and `double` use 8 bytes.

- In the previous code (`sum`), the memory accesses are made on `float` and the operational intensity is: $OI_{\text{sum}} = \frac{n \times 1}{(n \times 1) \times 4} = \frac{1}{4}$



Operational Intensity – Example

4 Kernel Performance Analysis

Sum in single precision:

```
1 // AI = 1  //  OI = 1/4
2 float sum1(float *values, size_t n) {
3     float sum = 0.f;
4     for (size_t i = 0; i < n; i++)
5         sum = sum + values[i];
6     return sum;
7 }
```

Sum in double precision:

```
1 // AI = 1  //  OI = 1/8
2 double sum2(double *values, size_t n) {
3     double sum = 0.0;
4     for (size_t i = 0; i < n; i++)
5         sum = sum + values[i];
6     return sum;
7 }
```

- `sum1` and `sum2` kernels have the same arithmetic intensity
- The operational intensity of `sum1` is higher than that of `sum2`
 - `sum2` kernel is more limited by memory access than `sum1` kernel



Roofline Model

4 Kernel Performance Analysis

- Roofline¹ is model that limits **the maximum achievable performance**
- Takes into account
 - Memory bandwidth (RAM)
 - Processor peak performance
- Depending on operational intensity, code is limited either by memory bandwidth or by processor peak performance

$$\text{Attainable op/s} = \min \begin{cases} \text{Peak CPU op/s performance,} \\ \text{Peak memory bandwidth} \times OI. \end{cases}$$

¹S. Williams, A. Waterman, and D. Patterson. “Roofline: An Insightful Visual Performance Model for Multicore Architectures”. In: *ACM Communications* 52.4 (Apr. 2009), pp. 65–76. DOI: 10.1145/1498765.1498785.



Memory Bandwidth Measurement

4 Kernel Performance Analysis

- Memory bandwidth or memory throughput represents the number of bytes that can be read/written from RAM in one second (o/s or GB/s)
- How to know the memory bandwidth?
 - It is possible to compute its theoretical value
 - But we often prefer to use a micro-benchmark program: “STREAM”¹ or “bandwidth”²
- STREAM (University of Virginia) is a small, relatively simple code for measuring memory bandwidth → C code
- bandwidth (Sorbonne University) targets the same features as STREAM but in a more friendly-user and accurate way → C++ code

¹STREAM: <https://www.cs.virginia.edu/stream/>

²bandwidth: <https://github.com/alsoc/bandwidth>



Roofline Model – Example – Part 1

4 Kernel Performance Analysis

Let's take the same processor as before:

CPU name	Core i7-2630QM
Peak perf sp	128 GFlop/s
Peak perf dp	64 GFlop/s
Mem. bandwidth	17.6 GB/s

Single precision sum:

```
1 // AI = 1  //  DI = 1/4
2 float sum1(float *values, size_t n) {
3     float sum = 0.f;
4     for (size_t i = 0; i < n; i++)
5         sum = sum + values[i];
6     return sum;
7 }
```

Double precision sum:

```
1 // AI = 1  //  DI = 1/8
2 double sum2(double *values, size_t n) {
3     double sum = 0.0;
4     for (size_t i = 0; i < n; i++)
5         sum = sum + values[i];
6     return sum;
7 }
```



Roofline Model – Example – Part 2

4 Kernel Performance Analysis

Peak perf sp	128 GFlop/s
Peak perf dp	64 GFlop/s
Mem. bandwidth	17.6 GB/s

For `sum2`, operational intensity is: $OI_{sum2} = 1/8$.

Let's apply the Roofline model:

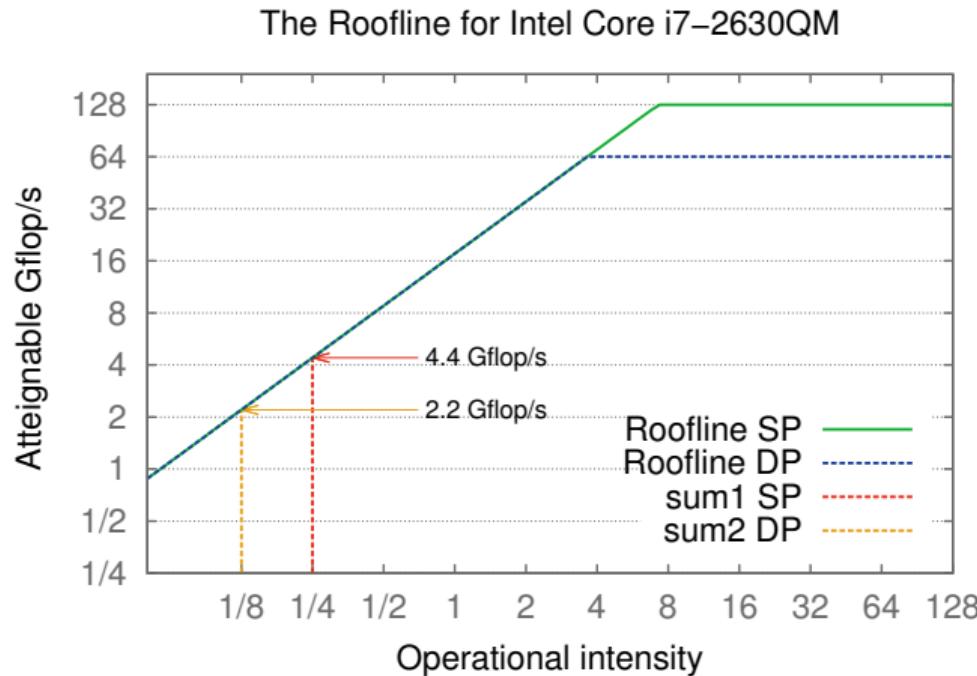
$$\text{Attainable Gflop/s} = \min \begin{cases} \text{Peak floating point performance,} \\ \text{Peak memory bandwidth} \times OI. \end{cases}$$

\Rightarrow

$$\text{Attainable Gflop/s}_{\text{sum2}} = \min \begin{cases} 64 \text{ Gflop/s,} \\ 17.6 \times \frac{1}{8} \text{ Gflop/s.} \end{cases} = 2.2 \text{ Gflop/s}$$

Roofline Model – Example – Visualization

4 Kernel Performance Analysis





Q&A

*Thank you for listening!
Do you have any questions?*



Bibliography

5 References

- [1] S. Williams, A. Waterman, and D. Patterson. “Roofline: An Insightful Visual Performance Model for Multicore Architectures”. In: *ACM Communications* 52.4 (Apr. 2009), pp. 65–76. DOI: [10.1145/1498765.1498785](https://doi.org/10.1145/1498765.1498785).